

# Final Project

Jessica Bao

3/15/2021

```
# Load libraries and prepare data
library("dplyr")
library("anytime")
library("lubridate")
library("readxl")
library("tidyverse")
library("TSA")
library("fpp")
library("Metrics")
library("corrplot")
library("tidymodels")

train <- read.csv("rossmann-store-sales/train.csv")
store_8 <- train %>% filter(Store==8) %>% arrange(Date)
#store_8

# Convert Date column to type "Date", get the first day of our data
store_8$Date = as.Date(store_8$Date, format = "%Y-%m-%d")

# Convert "03-01" to day of the year
dayOfYear = as.numeric(format(store_8[1,3], "%j"))

#keep the date information for future use
store_8$Date <- anydate(store_8$Date)
store_8$year <- year(store_8$Date)
store_8$month <- month(store_8$Date)
store_8$week <- week(store_8$Date)
store_8$day <- day(store_8$Date)

# double check 365 dates each year without missing
store_8 %>% count(year, sort = TRUE)

#take a look at the data: no missing data
#summary(store_8)

#drop store number
store_8 <- subset(store_8, select=-1)

# adjust the col sequence
store_8 <- subset(store_8, select= c(2:ncol(store_8),1))

#factor state holiday
```

```

store_8$StateHoliday<-factor(store_8$StateHoliday)
store_8$SchoolHoliday<-factor(store_8$SchoolHoliday)
store_8$DayOfWeek<-factor(store_8$DayOfWeek)
store_8$Open<-factor(store_8$Open)
store_8$Promo<-factor(store_8$Promo)

#convert to time series data
store_8 <- ts(store_8, frequency = 7)
#autoplot(store_8)

#split train test data
train <- window(store_8, start = c(1, 1), end = c(132, 4))
test <-window(store_8, start = c(132, 5))

#autoplot(train)
#autoplot(train[,2])

# Create a correlation matrix
train_corr <- cor(train)
round(train_corr, 2)[, 2] # Look at the correlations to Sales only

```

```

##      Date      Sales  Customers      Open      Promo
##      0.15      1.00      0.97      0.77      0.69
## StateHoliday SchoolHoliday      year      month      week
##      -0.25      0.11      0.13      0.04      0.04
##      day      DayOfWeek
##      -0.03      -0.71

```

*# Sales is weakly correlated with date, StateHoliday, SchoolHoliday, year, Month, Week, and Day.  
# Keep: Open, Promo, and DayofWeek.*

```

# Create the dependent and independent variables.
sales <- train[, "Sales"]
predictors <- train[, c("Open", "Promo", "DayOfWeek")]


```

```

# Differencing
sales_diff1 <- diff(sales, differences = 1)
tsdisplay(sales_diff1, main = "First Difference")

```

## First Difference



```
# ADF test
kpss.test(sales) #p = .01
```

```
## Warning in kpss.test(sales): p-value smaller than printed p-value
```

```
##
```

```
## KPSS Test for Level Stationarity
```

```
##
```

```
## data: sales
```

```
## KPSS Level = 2.8265, Truncation lag parameter = 6, p-value = 0.01
```

```
kpss.test(sales_diff1) #p = .1
```

```
## Warning in kpss.test(sales_diff1): p-value greater than printed p-value
```

```
##
```

```
## KPSS Test for Level Stationarity
```

```
##
```

```
## data: sales_diff1
```

```
## KPSS Level = 0.013299, Truncation lag parameter = 6, p-value = 0.1
```

The null hypothesis of the KPSS test is that the data is stationary. Because the p-value of the KPSS test of the first difference is less than alpha of .05, we reject the null, meaning the data is not stationary.

The p-value of the KPSS test of the second difference is greater than alpha of .05, so we fail to reject the null, meaning the data is stationary.

```
open <- diff(predictors[, "Open"], differences = 1)
promo <- diff(predictors[, "Promo"], differences = 1)
dayofweek <- diff(predictors[, "DayOfWeek"], differences = 1)
```

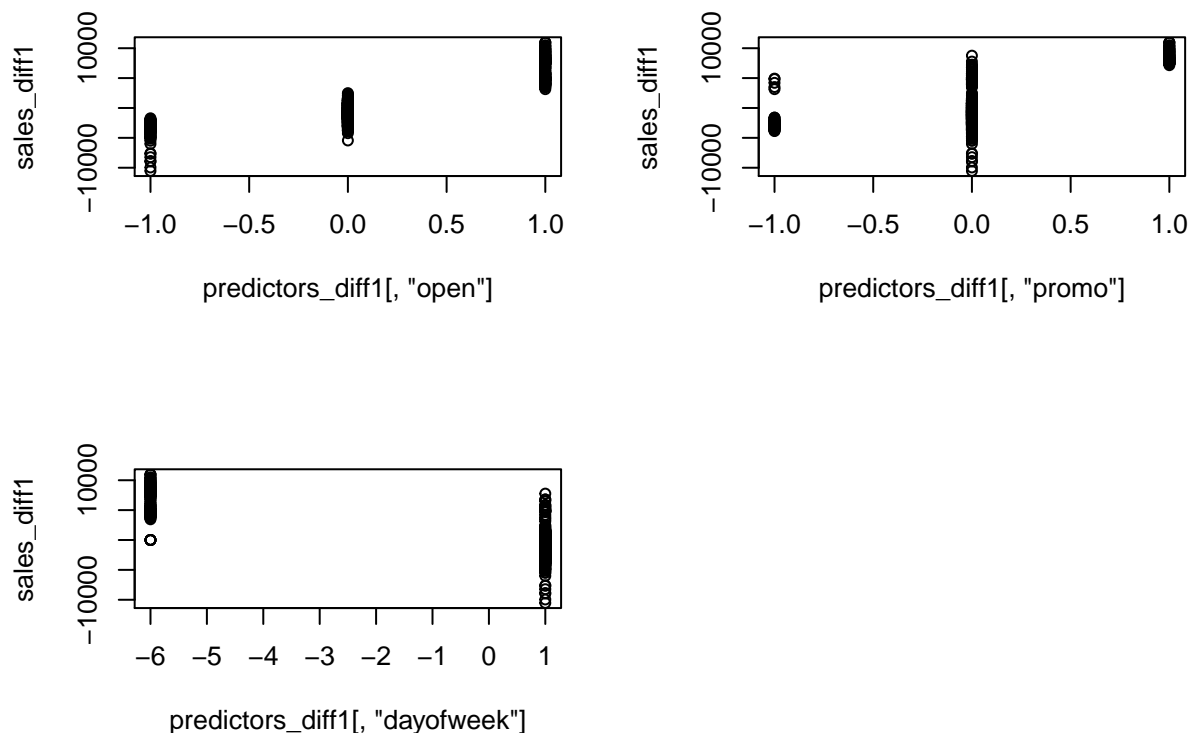
```
kpss.test(open) #p = .1
```

```
## Warning in kpss.test(open): p-value greater than printed p-value
##
## KPSS Test for Level Stationarity
##
## data: open
## KPSS Level = 0.04784, Truncation lag parameter = 6, p-value = 0.1
kpss.test(promo) #p = .1

## Warning in kpss.test(promo): p-value greater than printed p-value
##
## KPSS Test for Level Stationarity
##
## data: promo
## KPSS Level = 0.0045136, Truncation lag parameter = 6, p-value = 0.1
kpss.test(dayofweek) #p = .08411

##
## KPSS Test for Level Stationarity
##
## data: dayofweek
## KPSS Level = 0.38387, Truncation lag parameter = 6, p-value = 0.08411
predictors_diff1 = cbind(open, promo, dayofweek)

par(mfrow = c(2, 2))
plot(sales_diff1 ~ predictors_diff1[, "open"])
plot(sales_diff1 ~ predictors_diff1[, "promo"])
plot(sales_diff1 ~ predictors_diff1[, "dayofweek"])
```



```

# Fit a TSLM model.
tslm_fit <- tslm(sales_diff1 ~ predictors_diff1, lambda = "auto")

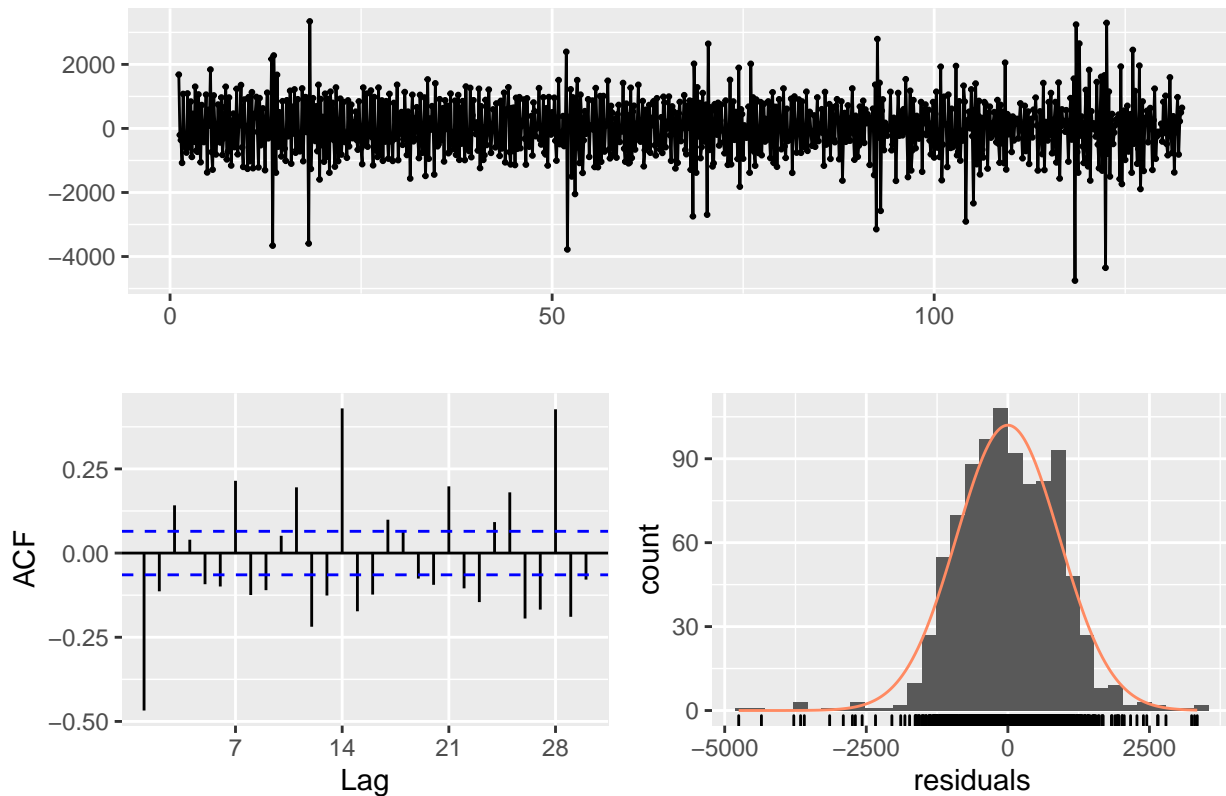
# Look at the summary of the model.
summary(tslm_fit)

##
## Call:
## tslm(formula = sales_diff1 ~ predictors_diff1, lambda = "auto")
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4752.6  -597.0   -21.0   651.7  3343.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -29.50      30.06  -0.981    0.327
## predictors_diff1open    2608.70      73.83  35.332 <2e-16 ***
## predictors_diff1promo    1978.67      93.40  21.185 <2e-16 ***
## predictors_diff1dayofweek -234.91      19.44 -12.082 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 911.9 on 916 degrees of freedom
## Multiple R-squared:  0.8688, Adjusted R-squared:  0.8684
## F-statistic: 2022 on 3 and 916 DF, p-value: < 2.2e-16

# Look at the residuals.
checkresiduals(tslm_fit)

```

## Residuals from Linear regression model



```
##
## Breusch-Godfrey test for serial correlation of order up to 14
##
## data: Residuals from Linear regression model
## LM test = 552.53, df = 14, p-value < 2.2e-16

# Look at AIC
#glance(tslm_fit)
AIC(tslm_fit, k = 3) # default k is 2 but because we want AICc, adding extra penalty parameter

## [1] 15162.39

# TSLM Box Ljung
Box.test(tslm_fit$residuals, type = c("Ljung-Box"))

##
## Box-Ljung test
##
## data: tslm_fit$residuals
## X-squared = 201.76, df = 1, p-value < 2.2e-16

# Fit regression with ARIMA errors
linear_arma_fit <- auto.arima(sales, xreg = predictors, lambda = "auto", seasonal = TRUE, allowdrift =

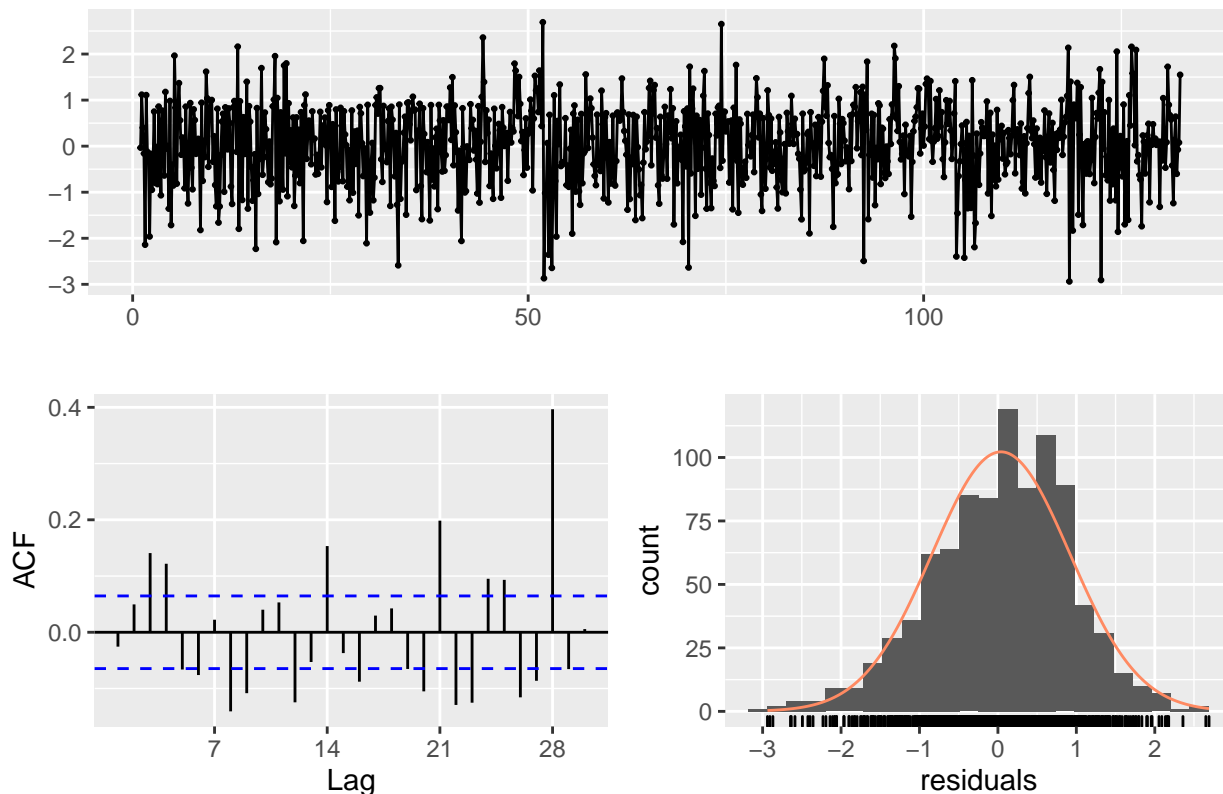
# Look at the summary of the model.
summary(linear_arma_fit)

## Series: sales
## Regression with ARIMA(1,1,2)(0,0,2)[7] errors
```

```
## Box Cox transformation: lambda= 0.2045416
##
## Coefficients:
##      ar1      ma1      ma2      sma1      sma2      Open      Promo      DayOfWeek
##      -0.9029 -0.0455 -0.9161  0.3345  0.4026  26.9716  2.3035  -0.3542
## s.e.    0.0512   0.0421   0.0408  0.0393  0.0293   0.1489   0.0647   0.0297
##
## sigma^2 estimated as 0.7824:  log likelihood=-1191.09
## AIC=2400.19   AICc=2400.38   BIC=2443.6
##
## Training set error measures:
##              ME      RMSE      MAE MPE MAPE      MASE      ACF1
## Training set 58.94127 740.8819 526.195 NaN  Inf 0.2935999 0.1634032

# Look at the residuals.
checkresiduals(linear_arima_fit)
```

### Residuals from Regression with ARIMA(1,1,2)(0,0,2)[7] errors



```
##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(1,1,2)(0,0,2)[7] errors
## Q* = 117.55, df = 6, p-value < 2.2e-16
##
## Model df: 8. Total lags used: 14

# Check EACF of sales
eacf(sales)
```

```
## AR/MA
```

```
##  0 1 2 3 4 5 6 7 8 9 10 11 12 13
## 0 x x x x x o x o x x x x x x x
## 1 x x o o x o x o x o o x o x
## 2 x x x o o x x x x x o o o x
## 3 x x o o o o x x o o o o o x
## 4 x x o o o x x x o o o o o x
## 5 x x o x x x o x x x o x o x
## 6 x x x x o x o x o o x x o x
## 7 x x x x o x x x o o x x x x
```

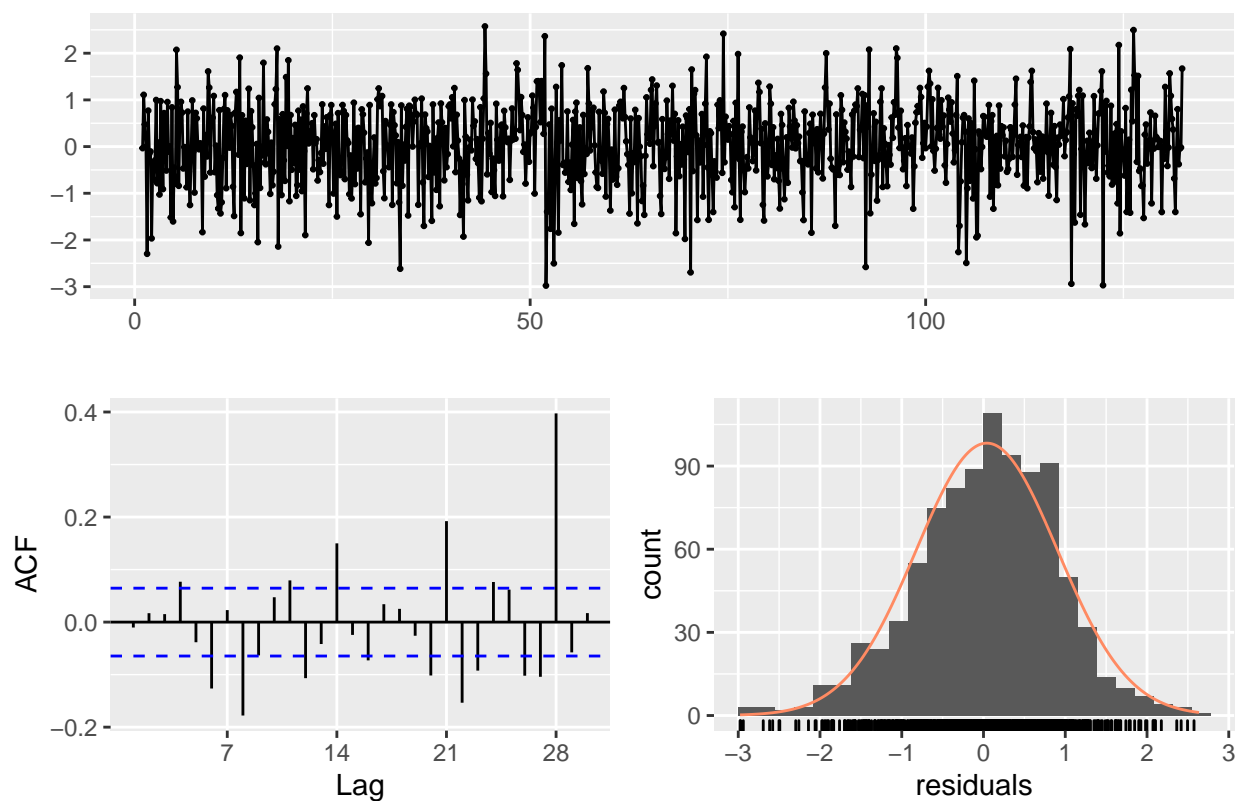
Try the following:

```
eacf1 <- Arima(sales, xreg = predictors, order = c(3, 1, 2), seasonal=list(order=c(0,0,2), period = 7),
summary(eacf1)
```

```
## Series: sales
## Regression with ARIMA(3,1,2)(0,0,2)[7] errors
## Box Cox transformation: lambda= 0.2045416
##
## Coefficients:
##          ar1      ar2      ar3      ma1      ma2      sma1      sma2      Open      Promo
##          0.2607  0.0412  0.1844 -1.2677  0.2771  0.3095  0.3982  26.9737  2.2855
## s.e.      0.0949  0.0352  0.0345   0.0939  0.0924  0.0396  0.0288   0.1341  0.0642
##          DayOfWeek
##          -0.3769
## s.e.      0.0255
##
## sigma^2 estimated as 0.7582:  log likelihood=-1175.57
## AIC=2373.14   AICc=2373.43   BIC=2426.21
##
## Training set error measures:
##              ME      RMSE      MAE MPE MAPE      MASE      ACF1
## Training set 60.93221 719.4383 511.3967 NaN  Inf  0.2853429 0.1843593
checkresiduals(eacf1)
```




# Residuals from Regression with ARIMA(3,1,2)(0,0,2)[7] errors



```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(3,1,2)(0,0,2)[7] errors
## Q* = 97.285, df = 4, p-value < 2.2e-16
##
## Model df: 10.   Total lags used: 14
acf(eacf1$residuals)
```

## Series eacf1\$residuals



```
# Split the test data into dependent and independent variables.
sales_test <- test[, "Sales"]
predictors_test <- test[, c("Open", "Promo", "DayOfWeek")]

# For TSLM because we're taking the differences, add the last train data
last_train <- train[921, 4]
last_predictors <- tail(predictors,1)

# Append to the test data
sales_test_for_tslm <- ts(c(last_train, sales_test), start = c(132, 4), frequency = frequency(sales_test))
predictors_test_for_tslm <- ts(rbind(last_predictors, predictors_test), start = c(132, 4), frequency = frequency(predictors_test))

# Take differences
sales_test_diff1 <- diff(sales_test_for_tslm, differences = 1)
open_test <- diff(predictors_test_for_tslm[, "Open"], differences = 1)
promo_test <- diff(predictors_test_for_tslm[, "Promo"], differences = 1)
dayofweek_test <- diff(predictors_test_for_tslm[, "DayOfWeek"], differences = 1)

predictors_test_diff1 = cbind(open_test, promo_test, dayofweek_test)
colnames(predictors_test_diff1) <- c("open", "promo", "dayofweek")

# Forecast the next three weeks
pred_tslm <- forecast(tslm_fit, h = 21, predictors_test_diff1, level = c(80, 95))

## Warning in forecast.lm(tslm_fit, h = 21, predictors_test_diff1, level = c(80, :
## newdata column names not specified, defaulting to first variable required.
pred_lm <- forecast(linear_arma_fit, h = 21, xreg = predictors_test, level= c(80, 95))

# Because the TSLM model used differenced data, its forecasts are also the differences between observat
```

```


# Take the last data point of the training data
pred_tslm_concat <- c(last_train, pred_tslm$mean)

# Add the TSLM predictions to the last observation and remove the first observation
pred_tslm_act <- cumsum(pred_tslm_concat)
pred_tslm_act <- head(pred_tslm_act, -1)

#ARIMA with linear error actuals + forecast
autoplot(pred_lm$mean, series = "Predicted", ylab = "Sales") + autolayer(sales_test, series = "Actuals")

```

sARIMA(1,1,2)(0,0,2)[7] Predicted vs Actual




```

#ARIMA with linear error ACF
acf(linear_arma_fit$residuals, main = "ACF of Residuals from sARIMA(1,1,2)(0,0,2)[7]")


```

### ACF of Residuals from sARIMA(1,1,2)(0,0,2)[7]

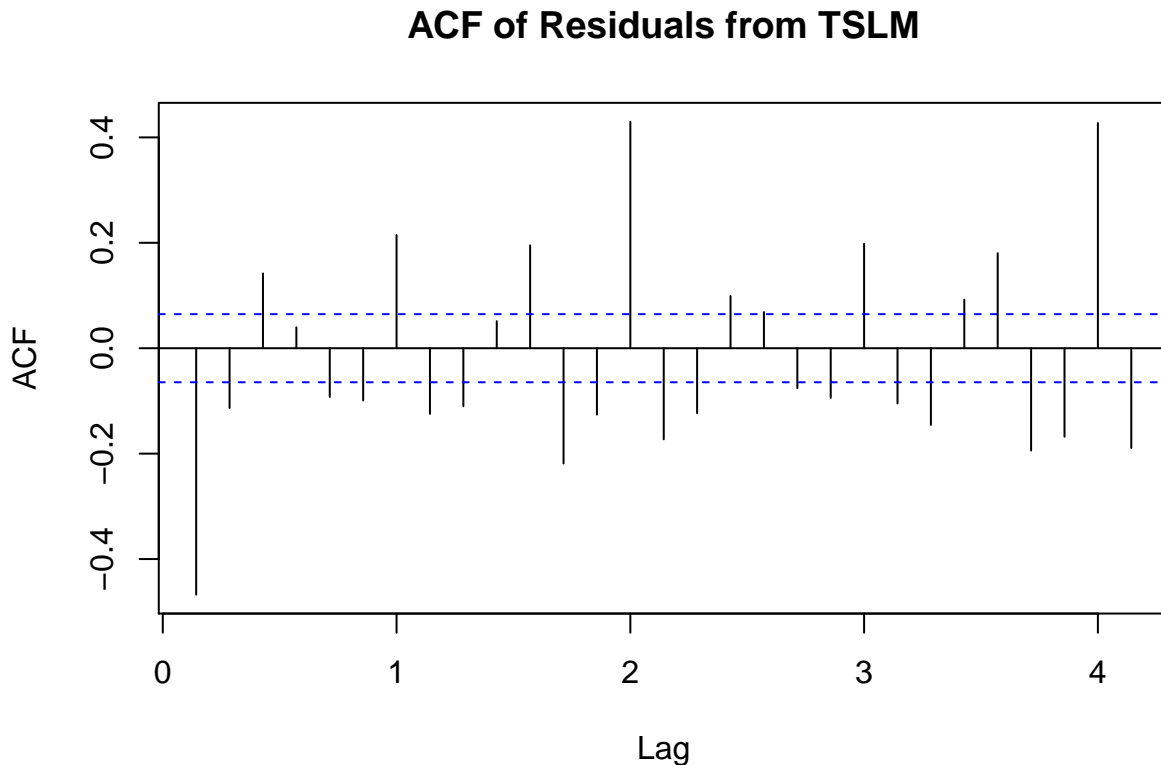


```
#TSLM of differences
autoplot(pred_tsml$mean, series = "Predicted") + autolayer(sales_test_diff1, series = "Actuals") + xlab
```

### TSLM Predicted vs Actual Differences



```
#TSLM ACF
acf(tslm_fit$residuals, main = "ACF of Residuals from TSLM")
```



```
# Convert test data to a dataframe
sales_test_df <- data.frame(Y=as.matrix(sales_test))

pred_lm_df <- data.frame(Y=as.matrix(pred_lm$mean))

lm_df <- cbind(sales_test_df, pred_lm_df)
colnames(lm_df) <- c("Y", "pred_lm_act")

tslm_df <- as.data.frame(cbind(sales_test_df, pred_tslm_act))
colnames(tslm_df) <- c("Y", "pred_tslm_act")

# RMSE
rmse_tslm <- rmse(tslm_df, "Y", "pred_tslm_act")
rmse_lm <- rmse(lm_df, "Y", "pred_lm_act")

cat("RMSE for TSLM model is ", rmse_tslm$.estimate)

## RMSE for TSLM model is 5748.325

cat("RMSE for linear model with ARIMA errors is ", rmse_lm$.estimate)

## RMSE for linear model with ARIMA errors is 764.3892

# Test RMSEs for the linear models with ARIMA error selected by EACF
pred_eacf1 <- forecast(eacf1, h = 21, xreg = predictors_test)
pred_eacf1_df <- data.frame(Y=as.matrix(pred_eacf1$mean))
eacf1_df <- cbind(sales_test_df, pred_eacf1_df)
colnames(eacf1_df) <- c("Y", "pred_eacf1_act")
```

```

rmse_eacf1 <- rmse(eacf1_df, "Y", "pred_eacf1_act")

cat("RMSE for ARIMA with sARIMA(3,1,2)(0,0,7) model is ", rmse_eacf1$.estimate)

## RMSE for ARIMA with sARIMA(3,1,2)(0,0,7) model is 790.1098

ADDING IN ALL PREDICTORS

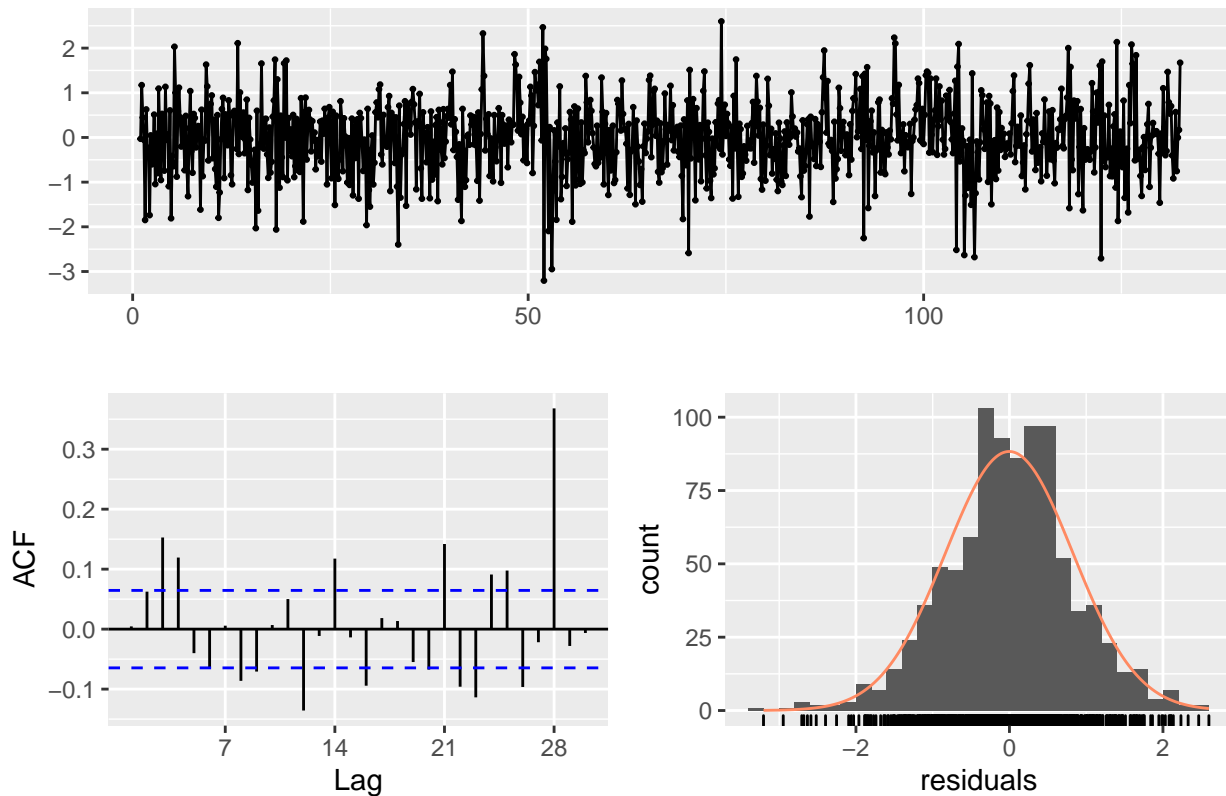
# Retrain the model using all predictors
predictors_all <- train[, c("Open", "Promo", "StateHoliday", "SchoolHoliday", "DayOfWeek")] #excluded:
linear_arma_fit_all <- auto.arima(sales, xreg = predictors_all, lambda = "auto", seasonal = TRUE, allow
summary(linear_arma_fit_all)

## Series: sales
## Regression with ARIMA(0,1,2)(0,0,2)[7] errors
## Box Cox transformation: lambda= 0.2045416
##
## Coefficients:
##          ma1      ma2      sma1      sma2      drift      Open      Promo      StateHoliday
##      -0.8801  -0.1016   0.3129   0.3478   0.0016  25.7623   2.3166        -1.3564
## s.e.    0.0356   0.0362   0.0377   0.0285   0.0009   0.1608   0.0654         0.1331
##      SchoolHoliday DayOfWeek
##              0.1365    -0.4723
## s.e.          0.1010     0.0297
##
## sigma^2 estimated as 0.7008: log likelihood=-1139.14
## AIC=2300.29  AICc=2300.58  BIC=2353.35
##
## Training set error measures:
##              ME      RMSE      MAE MPE MAPE      MASE      ACF1
## Training set 38.88672 718.8552 509.9592 NaN  Inf  0.2845409 0.1571306

checkresiduals(linear_arma_fit_all) #Ljung Box p-value is 2.2e-16

```

## Residuals from Regression with ARIMA(0,1,2)(0,0,2)[7] errors



```
##
##  Ljung-Box test
##
## data:  Residuals from Regression with ARIMA(0,1,2)(0,0,2)[7] errors
## Q* = 88.181, df = 4, p-value < 2.2e-16
##
## Model df: 10.   Total lags used: 14

# Reforecast using updated model
predictors_test_all <- test[, c("Open", "Promo", "StateHoliday", "SchoolHoliday", "DayOfWeek")]
pred_lm_all <- forecast(linear_arma_fit_all, h = 21, xreg = predictors_test_all, level= c(80, 95))

# Reformat the predicted data
pred_lm_df_all <- data.frame(Y=as.matrix(pred_lm_all$mean))
lm_df_all <- cbind(sales_test_df, pred_lm_df_all)
colnames(lm_df_all) <- c("Y", "pred_lm_act_all")


# RMSE
rmse_lm_all <- rmse(lm_df_all, "Y", "pred_lm_act_all")
cat("RMSE for ARIMA with linear errors model is ", rmse_lm_all$estimate)

## RMSE for ARIMA with linear errors model is  830.5224

#ARIMA with linear error actuals + forecast
pred_eacf1 <- forecast(eacf1, h = 21, xreg = predictors_test, level= c(80, 95))


autoplot(pred_eacf1$mean, series = "Predicted", ylab = "Sales") + autolayer(sales_test, series = "Actual")
```

sARIMA(3,1,2)(0,0,2)[7] Predicted vs Actual



```
#EACF ARIMA with linear error ACF
acf(eacf1$residuals, main = "ACF of Residuals from sARIMA(3,1,2)(0,0,2)[7]")
```

ACF of Residuals from sARIMA(3,1,2)(0,0,2)[7]

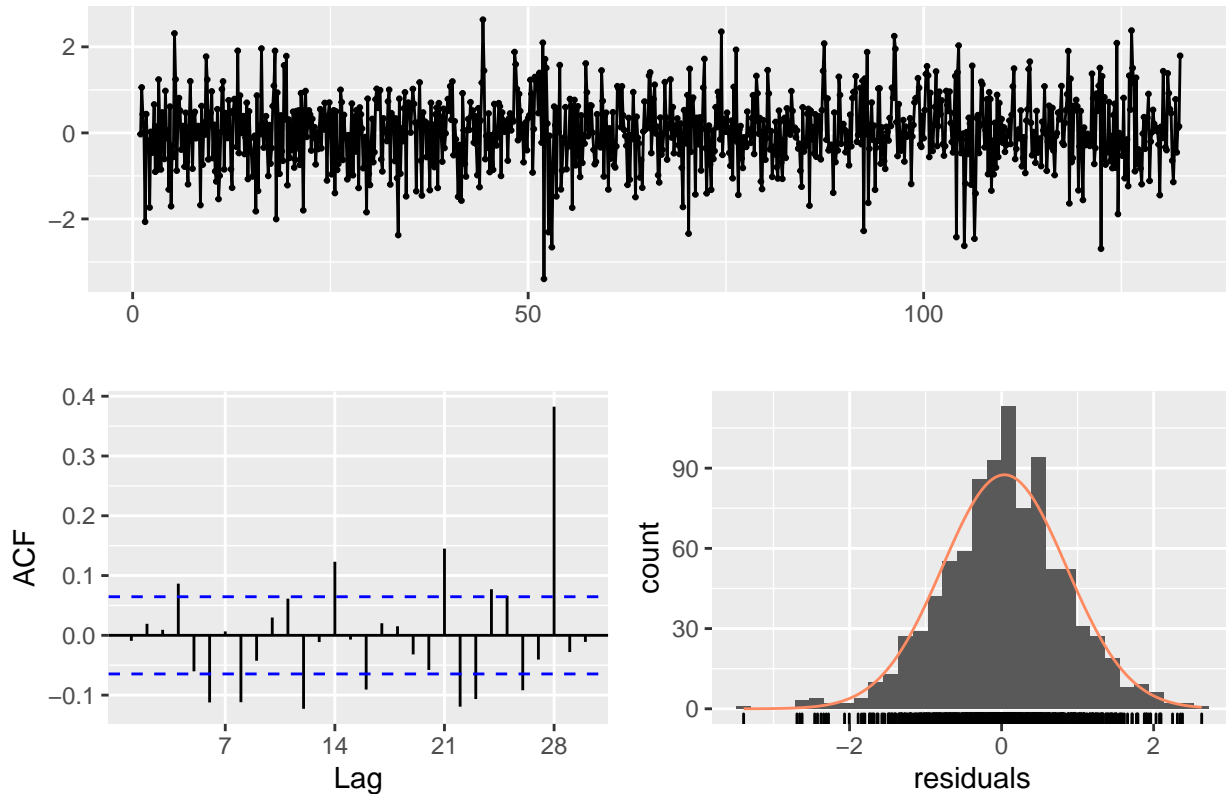




```
# Retrain using all predictors
```

```
eacf1_all <- Arima(sales, xreg = predictors_all, order = c(3, 1, 2), seasonal=list(order=c(0,0,2), period=12))
checkresiduals(eacf1_all)
```

Residuals from Regression with ARIMA(3,1,2)(0,0,2)[7] errors



```
##
```

```
## Ljung-Box test
```

```
##
```

```
## data: Residuals from Regression with ARIMA(3,1,2)(0,0,2)[7] errors
```

```
## Q* = 68.693, df = 3, p-value = 8.105e-15
```

```
##
```

```
## Model df: 12. Total lags used: 15
```

```
# Predict
```

```
pred_eacf1_all <- forecast(eacf1_all, h = 21, xreg = predictors_test_all, level= c(80, 95))
```

```
# Reformat the predicted data
```

```
pred_eacf1_df_all <- data.frame(Y=as.matrix(pred_eacf1_all$mean))
```

```
eacf_df_all <- cbind(sales_test_df, pred_eacf1_df_all)
```

```
colnames(eacf_df_all) <- c("Y", "pred_lm_act_all")
```

```
# RMSE
```

```
rmse_eacf1_all <- rmse(eacf_df_all, "Y", "pred_lm_act_all")
```

```
cat("RMSE for ARIMA with linear errors model is ", rmse_eacf1_all$estimate)
```

```
## RMSE for ARIMA with linear errors model is 847.2358
```